

2019-2024-market-data-analysis

March 29, 2025

1 Analysis Project

The goal here is to exercise the use of data cleaning and data visualization techniques for the use of data analytics.

The project focuses on US stock market data.

The steps taken in this exercise: 1. Review Metadata 2. Import Kaggle Dataset 3. Data Exploration & Cleaning 4. Organize Data Into Groups 5. Stock Visualizations 6. Market Visualizations 7. Commodity Visualizations 8. Combined Analysis

1.0.1 1) Reviewing Metadata

It is critical to gain a contextual understanding before touching the data by reviewing the metadata.

- **Date:** The date of the recorded data, formatted as DD-MM-YYYY.
- **Natural_Gas_Price:** Price of natural gas in USD per million British thermal units (MMBtu).
- **Natural_Gas_Vol.:** Trading volume of natural gas.
- **Crude_oil_Price:** Price of crude oil in USD per barrel.
- **Crude_oil_Vol.:** Trading volume of crude oil.
- **Copper_Price:** Price of copper in USD per pound.
- **Copper_Vol.:** Trading volume of copper.
- **Bitcoin_Price:** Price of Bitcoin in USD.
- **Bitcoin_Vol.:** Trading volume of Bitcoin.
- **Platinum_Price:** Price of platinum in USD per troy ounce.
- **Platinum_Vol.:** Trading volume of platinum.
- **Ethereum_Price:** Price of Ethereum in USD.
- **Ethereum_Vol.:** Trading volume of Ethereum.
- **S&P_500_Price:** Price index of the S&P 500.
- **Nasdaq_100_Price:** Price index of the Nasdaq 100.
- **Nasdaq_100_Vol.:** Trading volume for the Nasdaq 100 index.
- **Apple_Price:** Stock price of Apple Inc. in USD.
- **Apple_Vol.:** Trading volume of Apple Inc. stock.
- **Tesla_Price:** Stock price of Tesla Inc. in USD.
- **Tesla_Vol.:** Trading volume of Tesla Inc. stock.
- **Microsoft_Price:** Stock price of Microsoft Corporation in USD.
- **Microsoft_Vol.:** Trading volume of Microsoft Corporation stock.
- **Silver_Price:** Price of silver in USD per troy ounce.
- **Silver_Vol.:** Trading volume of silver.

- **Google_Price**: Stock price of Alphabet Inc. (Google) in USD.
- **Google_Vol.**: Trading volume of Alphabet Inc. stock.
- **Nvidia_Price**: Stock price of Nvidia Corporation in USD.
- **Nvidia_Vol.**: Trading volume of Nvidia Corporation stock.
- **Berkshire_Price**: Stock price of Berkshire Hathaway Inc. in USD.
- **Berkshire_Vol.**: Trading volume of Berkshire Hathaway Inc. stock.
- **Netflix_Price**: Stock price of Netflix Inc. in USD.
- **Netflix_Vol.**: Trading volume of Netflix Inc. stock.
- **Amazon_Price**: Stock price of Amazon.com Inc. in USD.
- **Amazon_Vol.**: Trading volume of Amazon.com Inc. stock.
- **Meta_Price**: Stock price of Meta Platforms, Inc. (formerly Facebook) in USD.
- **Meta_Vol.**: Trading volume of Meta Platforms, Inc. stock.
- **Gold_Price**: Price of gold in USD per troy ounce.
- **Gold_Vol.**: Trading volume of gold.

Reference: <https://www.kaggle.com/datasets/saketk511/2019-2024-us-stock-market-data>.

1.0.2 2) Importing Dataset

Gather the stock information data from a source in Kaggle. The file is taken from a directory and converted into a dataframe on the notebook, where it'll be analyzed.

```
[1]: import kagglehub

# Download latest version
data_path = kagglehub.dataset_download("saketk511/
↳2019-2024-us-stock-market-data")

print("Path to dataset files:", data_path)
```

Downloading from

https://www.kaggle.com/api/v1/datasets/download/saketk511/2019-2024-us-stock-market-data?dataset_version_number=1...

100%| | 155k/155k [00:00<00:00, 1.00MB/s]

Extracting files...

Path to dataset files:

/Users/chevalier/.cache/kagglehub/datasets/saketk511/2019-2024-us-stock-market-data/versions/1

```
[2]: import os

files = os.listdir(data_path)
print(files)
```

['Stock Market Dataset.csv']

```
[3]: file_path = os.path.join(data_path, files[0])
      print(file_path)
```

/Users/chevalier/.cache/kagglehub/datasets/saketk511/2019-2024-us-stock-market-data/versions/1/Stock Market Dataset.csv

```
[4]: import pandas as pd

df = pd.read_csv(file_path, index_col=0)
df.head(5)
```

```
[4]:
```

	Date	Natural_Gas_Price	Natural_Gas_Vol.	Crude_oil_Price	\
0	02-02-2024	2.079	NaN	72.28	
1	01-02-2024	2.050	161340.0	73.82	
2	31-01-2024	2.100	142860.0	75.85	
3	30-01-2024	2.077	139750.0	77.82	
4	29-01-2024	2.490	3590.0	76.78	

	Crude_oil_Vol.	Copper_Price	Copper_Vol.	Bitcoin_Price	Bitcoin_Vol.	\
0	NaN	3.8215	NaN	43,194.70	42650.0	
1	577940.0	3.8535	NaN	43,081.40	47690.0	
2	344490.0	3.9060	NaN	42,580.50	56480.0	
3	347240.0	3.9110	NaN	42,946.20	55130.0	
4	331930.0	3.8790	NaN	43,299.80	45230.0	

	Platinum_Price	...	Berkshire_Price	Berkshire_Vol.	Netflix_Price	\
0	901.6	...	5,89,498	10580.0	564.64	
1	922.3	...	5,81,600	9780.0	567.51	
2	932.6	...	5,78,020	9720.0	564.11	
3	931.7	...	5,84,680	9750.0	562.85	
4	938.3	...	5,78,800	13850.0	575.79	

	Netflix_Vol.	Amazon_Price	Amazon_Vol.	Meta_Price	Meta_Vol.	Gold_Price	\
0	4030000.0	171.81	117220000.0	474.99	84710000.0	2,053.70	
1	3150000.0	159.28	66360000.0	394.78	25140000.0	2,071.10	
2	4830000.0	155.20	49690000.0	390.14	20010000.0	2,067.40	
3	6120000.0	159.00	42290000.0	400.06	18610000.0	2,050.90	
4	6880000.0	161.26	42840000.0	401.02	17790000.0	2,034.90	

	Gold_Vol.
0	NaN
1	260920.0
2	238370.0
3	214590.0
4	1780.0

[5 rows x 38 columns]

1.0.3 3) Data Exploration & Cleaning

Understand the data and prepare it for future analysis.

```
[5]: df.shape
```

```
[5]: (1243, 38)
```

Inspect how many days were not captured for each of the columns.

```
[6]: nulls=df.isnull().sum().to_frame()  
nulls.rename(columns={0:'Nulls'},inplace=True)  
nulls
```

```
[6]:
```

	Nulls
Date	0
Natural_Gas_Price	0
Natural_Gas_Vol.	4
Crude_oil_Price	0
Crude_oil_Vol.	23
Copper_Price	0
Copper_Vol.	37
Bitcoin_Price	0
Bitcoin_Vol.	0
Platinum_Price	0
Platinum_Vol.	607
Ethereum_Price	0
Ethereum_Vol.	0
S&P_500_Price	0
Nasdaq_100_Price	0
Nasdaq_100_Vol.	1
Apple_Price	0
Apple_Vol.	0
Tesla_Price	0
Tesla_Vol.	0
Microsoft_Price	0
Microsoft_Vol.	0
Silver_Price	0
Silver_Vol.	47
Google_Price	0
Google_Vol.	0
Nvidia_Price	0
Nvidia_Vol.	0
Berkshire_Price	0
Berkshire_Vol.	0
Netflix_Price	0
Netflix_Vol.	0
Amazon_Price	0
Amazon_Vol.	0

Meta_Price	0
Meta_Vol.	0
Gold_Price	0
Gold_Vol.	2

```
[7]: columns_to_fill = ['Natural_Gas_Vol.', 'Crude_oil_Vol.', 'Copper_Vol.',
↳ ', 'Platinum_Vol.', 'Nasdaq_100_Vol.', 'Silver_Vol.', 'Gold_Vol.']

for column in columns_to_fill:
    df[column] = df[column].ffill()
    df[column] = df[column].bfill()
```

```
[8]: nulls=df.isnull().sum().to_frame()
nulls.rename(columns={0: 'Nulls'}, inplace=True)
nulls
```

```
[8]:
```

	Nulls
Date	0
Natural_Gas_Price	0
Natural_Gas_Vol.	0
Crude_oil_Price	0
Crude_oil_Vol.	0
Copper_Price	0
Copper_Vol.	0
Bitcoin_Price	0
Bitcoin_Vol.	0
Platinum_Price	0
Platinum_Vol.	0
Ethereum_Price	0
Ethereum_Vol.	0
S&P_500_Price	0
Nasdaq_100_Price	0
Nasdaq_100_Vol.	0
Apple_Price	0
Apple_Vol.	0
Tesla_Price	0
Tesla_Vol.	0
Microsoft_Price	0
Microsoft_Vol.	0
Silver_Price	0
Silver_Vol.	0
Google_Price	0
Google_Vol.	0
Nvidia_Price	0
Nvidia_Vol.	0
Berkshire_Price	0
Berkshire_Vol.	0

Netflix_Price	0
Netflix_Vol.	0
Amazon_Price	0
Amazon_Vol.	0
Meta_Price	0
Meta_Vol.	0
Gold_Price	0
Gold_Vol.	0

Verify that there are no duplicated days, which could potentially hinder analysis.

```
[9]: df[df.duplicated()]
```

```
[9]: Empty DataFrame
Columns: [Date, Natural_Gas_Price, Natural_Gas_Vol., Crude_oil_Price,
Crude_oil_Vol., Copper_Price, Copper_Vol., Bitcoin_Price, Bitcoin_Vol.,
Platinum_Price, Platinum_Vol., Ethereum_Price, Ethereum_Vol., S&P_500_Price,
Nasdaq_100_Price, Nasdaq_100_Vol., Apple_Price, Apple_Vol., Tesla_Price,
Tesla_Vol., Microsoft_Price, Microsoft_Vol., Silver_Price, Silver_Vol.,
Google_Price, Google_Vol., Nvidia_Price, Nvidia_Vol., Berkshire_Price,
Berkshire_Vol., Netflix_Price, Netflix_Vol., Amazon_Price, Amazon_Vol.,
Meta_Price, Meta_Vol., Gold_Price, Gold_Vol.]
Index: []
```

[0 rows x 38 columns]

Change the data types of date and the other columns to be datetime and float formats.

```
[10]: df['Date']=pd.to_datetime(df.Date, format='%d-%m-%Y')

for col in df.columns[1:]:
    df[col] = df[col].replace(',', '', regex=True).astype('float64')
```

Verify the data types are numerical for all columns excluding the date.

```
[11]: types=df.dtypes.to_frame()
types.rename(columns={0:'Types'},inplace=True)
types
```

```
[11]:
```

	Types
Date	datetime64[ns]
Natural_Gas_Price	float64
Natural_Gas_Vol.	float64
Crude_oil_Price	float64
Crude_oil_Vol.	float64
Copper_Price	float64
Copper_Vol.	float64
Bitcoin_Price	float64
Bitcoin_Vol.	float64

Platinum_Price	float64
Platinum_Vol.	float64
Ethereum_Price	float64
Ethereum_Vol.	float64
S&P_500_Price	float64
Nasdaq_100_Price	float64
Nasdaq_100_Vol.	float64
Apple_Price	float64
Apple_Vol.	float64
Tesla_Price	float64
Tesla_Vol.	float64
Microsoft_Price	float64
Microsoft_Vol.	float64
Silver_Price	float64
Silver_Vol.	float64
Google_Price	float64
Google_Vol.	float64
Nvidia_Price	float64
Nvidia_Vol.	float64
Berkshire_Price	float64
Berkshire_Vol.	float64
Netflix_Price	float64
Netflix_Vol.	float64
Amazon_Price	float64
Amazon_Vol.	float64
Meta_Price	float64
Meta_Vol.	float64
Gold_Price	float64
Gold_Vol.	float64

Inspect some descriptive statistics of the dataset.

```
[12]: df.describe()
```

```
[12]:
```

	Date	Natural_Gas_Price	Natural_Gas_Vol.	\
count	1243	1243.000000	1243.000000	
mean	2021-08-02 10:03:34.320193024	3.494714	131572.510056	
min	2019-02-04 00:00:00	1.482000	1200.000000	
25%	2020-04-28 12:00:00	2.347500	91700.000000	
50%	2021-08-03 00:00:00	2.702000	127370.000000	
75%	2022-11-05 12:00:00	4.055500	169340.000000	
max	2024-02-02 00:00:00	9.647000	381970.000000	
std	NaN	1.822540	64320.925408	

	Crude_oil_Price	Crude_oil_Vol.	Copper_Price	Copper_Vol.	\
count	1243.000000	1.243000e+03	1243.000000	1243.000000	
mean	67.577064	3.944626e+05	3.541957	35668.270314	
min	-37.630000	1.702000e+04	2.100500	10.000000	

25%	55.095000	2.774400e+05	2.858750	380.000000
50%	69.230000	3.647200e+05	3.666000	21530.000000
75%	80.455000	5.037300e+05	4.137250	68030.000000
max	123.700000	1.770000e+06	4.937500	176040.000000
std	20.465500	2.169229e+05	0.702819	38006.113950

	Bitcoin_Price	Bitcoin_Vol.	Platinum_Price	...	Berkshire_Price	\
count	1243.000000	1.243000e+03	1243.000000	...	1243.000000	
mean	25241.903057	4.033918e+07	959.003620	...	404273.051488	
min	3397.700000	2.600000e+02	595.200000	...	240000.000000	
25%	10014.600000	7.907500e+04	889.775000	...	318984.500000	
50%	23055.100000	2.153100e+05	944.700000	...	418349.000000	
75%	37784.200000	6.151050e+05	1020.400000	...	471500.000000	
max	67527.900000	4.470000e+09	1297.100000	...	589498.000000	
std	16029.009055	2.940889e+08	108.012849	...	86369.903899	

	Berkshire_Vol.	Netflix_Price	Netflix_Vol.	Amazon_Price	\
count	1243.000000	1243.000000	1.243000e+03	1243.000000	
mean	2426.524537	404.839541	7.057401e+06	128.683234	
min	80.000000	166.370000	1.140000e+06	79.410000	
25%	345.000000	323.010000	3.990000e+06	96.260000	
50%	1510.000000	384.150000	5.610000e+06	128.730000	
75%	3225.000000	495.365000	7.910000e+06	158.110000	
max	13850.000000	691.690000	1.333900e+08	186.570000	
std	2660.497572	114.989473	6.384187e+06	30.808631	

	Amazon_Vol.	Meta_Price	Meta_Vol.	Gold_Price	Gold_Vol.
count	1.243000e+03	1243.000000	1.243000e+03	1243.000000	1243.000000
mean	7.413005e+07	239.728134	2.325851e+07	1759.246742	210998.238134
min	1.763000e+07	88.910000	5.470000e+06	1272.000000	0.000000
25%	5.264500e+07	183.355000	1.478500e+07	1669.600000	152060.000000
50%	6.520000e+07	224.430000	1.934000e+07	1804.200000	197970.000000
75%	8.674500e+07	301.650000	2.711500e+07	1912.800000	258370.000000
max	3.113500e+08	474.990000	2.304100e+08	2089.700000	813410.000000
std	3.245753e+07	71.015427	1.555486e+07	203.258901	115077.705917

[8 rows x 38 columns]

Key Learnings The timeframe captured by the dataset spans between February 2019 and February 2024.

There are 1243 days of information gathered. For each day, we have prices and volumes of commodities, crypto, markets, and stocks.

1.0.4 4) Organize Data Into Groups

Separate the data so it is comparable to other similar columns.


```
[13]: stock_prices = df[['Apple_Price', 'Tesla_Price', 'Microsoft_Price',
    ↪ 'Google_Price', 'Nvidia_Price', 'Berkshire_Price',
    ↪ 'Netflix_Price', 'Amazon_Price', 'Meta_Price']]

market_prices = df[['S&P_500_Price', 'Nasdaq_100_Price']]

commodity_prices = df[['Natural_Gas_Price', 'Crude_oil_Price', 'Copper_Price',
    ↪ 'Platinum_Price', 'Silver_Price', 'Gold_Price']]
```

1.0.5 5) Stock Visualizations

```
[14]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

Line Plots:

Analyzing the movements of key stocks over time.

Companies with the highest growth since 2019 include Nvidia and Microsoft.

Netflix grew the least of the bunch.


```
[15]: colors = sns.color_palette("tab10", len(stock_prices.columns))

fig, axes = plt.subplots(nrows=len(stock_prices.columns), ncols=1,
    ↪ figsize=(12,10), sharex=True)

for i, (stock,color) in enumerate(zip(stock_prices.columns,colors)):
    axes[i].plot(df['Date'], df[stock], label=stock, color=color)
    axes[i].set_title(f'{stock} Over Time')
    axes[i].set_ylabel('Price (USD)')
    axes[i].legend()

plt.xticks(rotation=45)
plt.xlabel('Date')


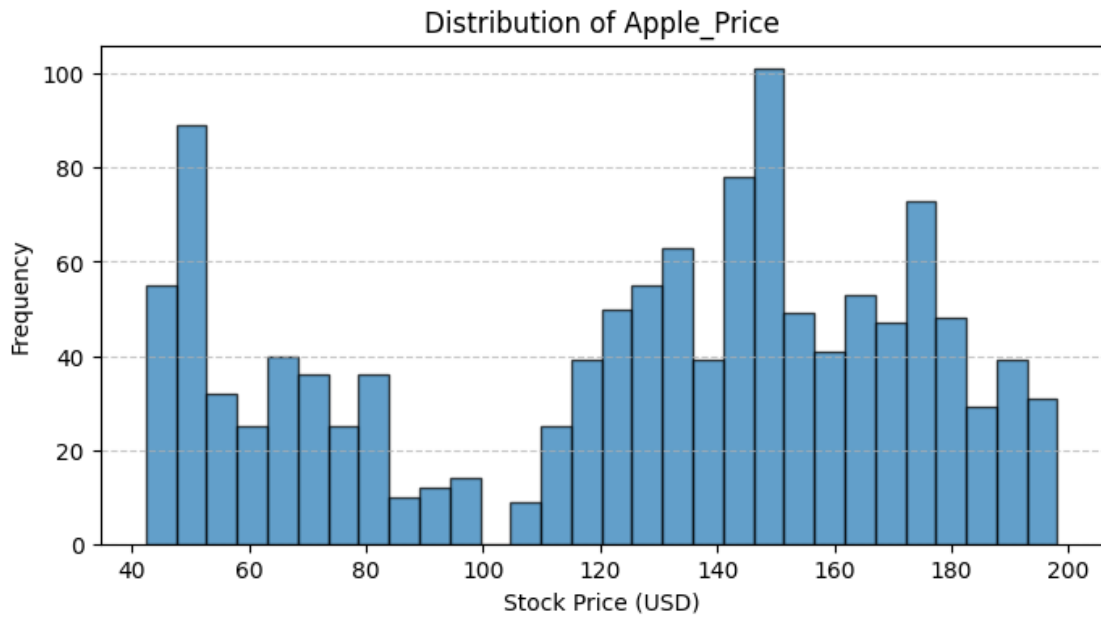
plt.tight_layout()
plt.show()
```





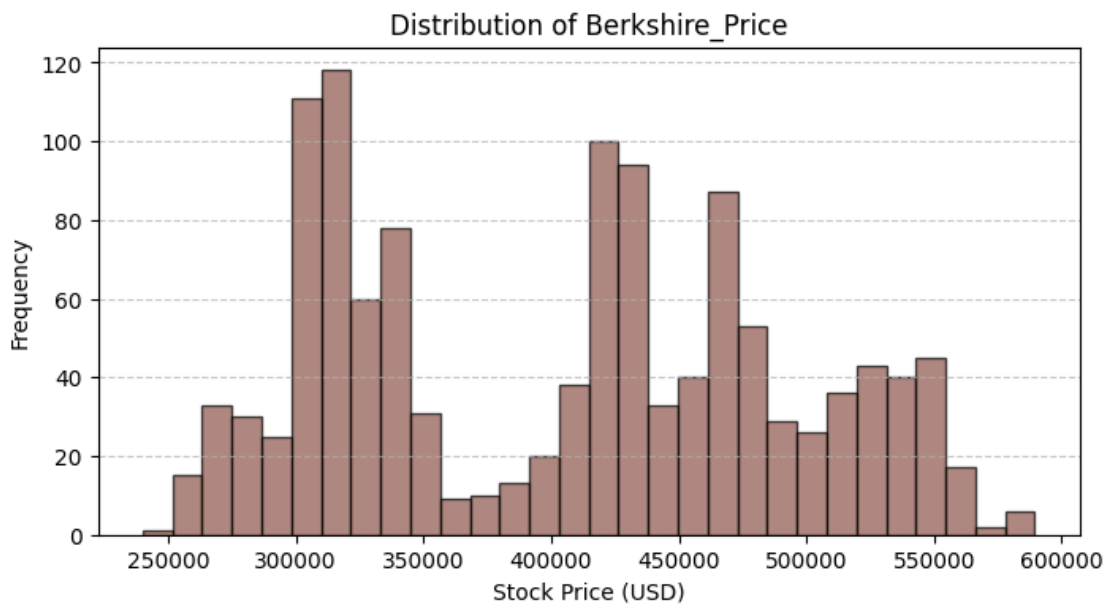

Histogram:

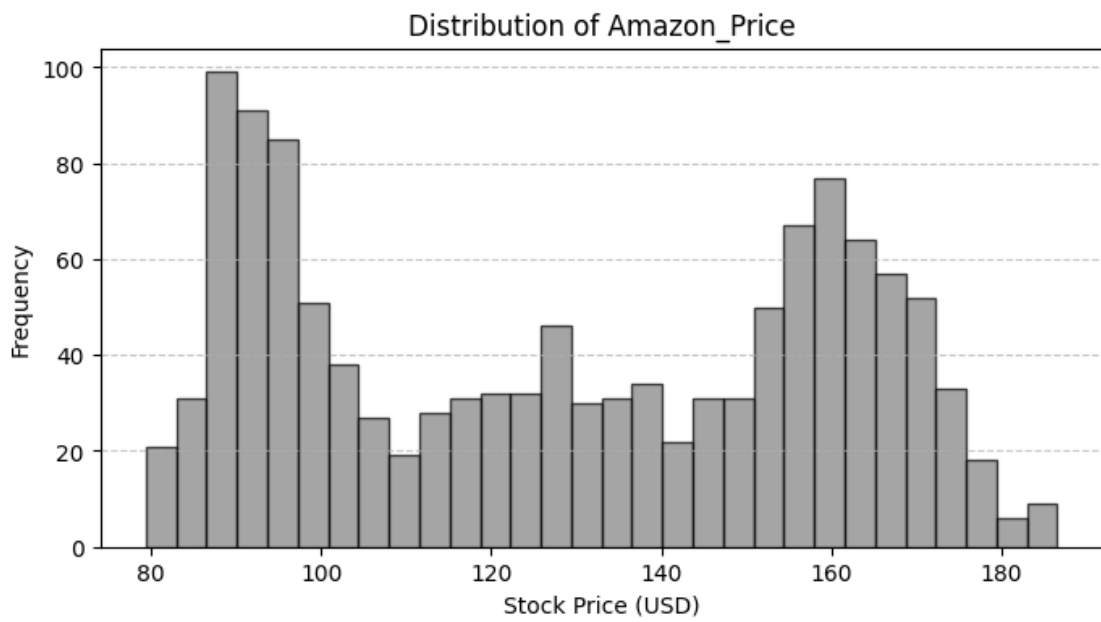

Check distribution of prices over the dataset period.

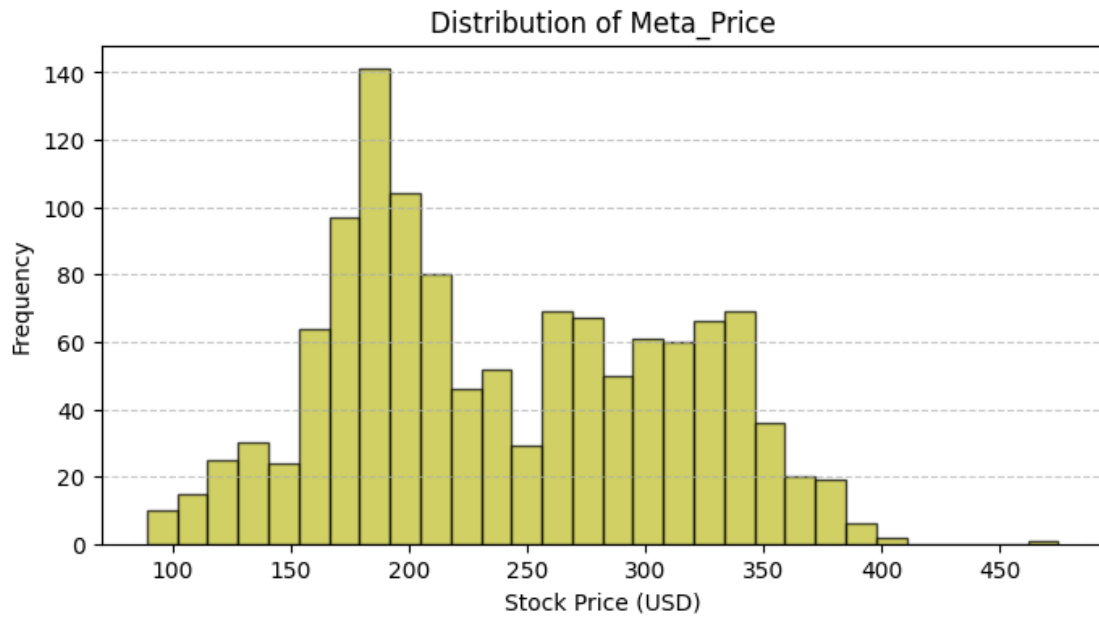
```
[16]: for stock, color in zip(stock_prices.columns, colors):
    plt.figure(figsize=(8, 4))
    plt.hist(stock_prices[stock], bins=30, color=color, alpha=0.7,
edgecolor='black')
    plt.title(f'Distribution of {stock}')
    plt.xlabel('Stock Price (USD)')
    plt.ylabel('Frequency')
    plt.grid(axis='y', linestyle='--', alpha=0.7)
    plt.show()
```












Heat Map:

Check for correlations of stock prices between companies.

```
[17]: corr_matrix = stock_prices.corr()

plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title('Stock Price Correlation Heatmap')
plt.show()
```



Microsoft is the most correlated to the other stocks analyzed (old/new Mag7 and Berkshire).

Netflix is the least connected with a correlation significantly lower than the others. This is logical as it is a streaming service, while the others are mainly technology oriented (Berkshire being an exception).

Berkshire has investments in Apple (40-50% of its portfolio) and Amazon. This is why Berkshire, despite not being a tech company, has a strong positive correlation 0.9.

```
[18]: average_correlations=[]
      for column in corr_matrix.columns:
          average_correlations.append([column,corr_matrix[column].mean()])
ac_frame = pd.DataFrame(average_correlations,columns=['Stocks','Average_
↳Correlation']).sort_values(by='Average Correlation',ascending=False).
↳set_index('Stocks')
ac_frame
```

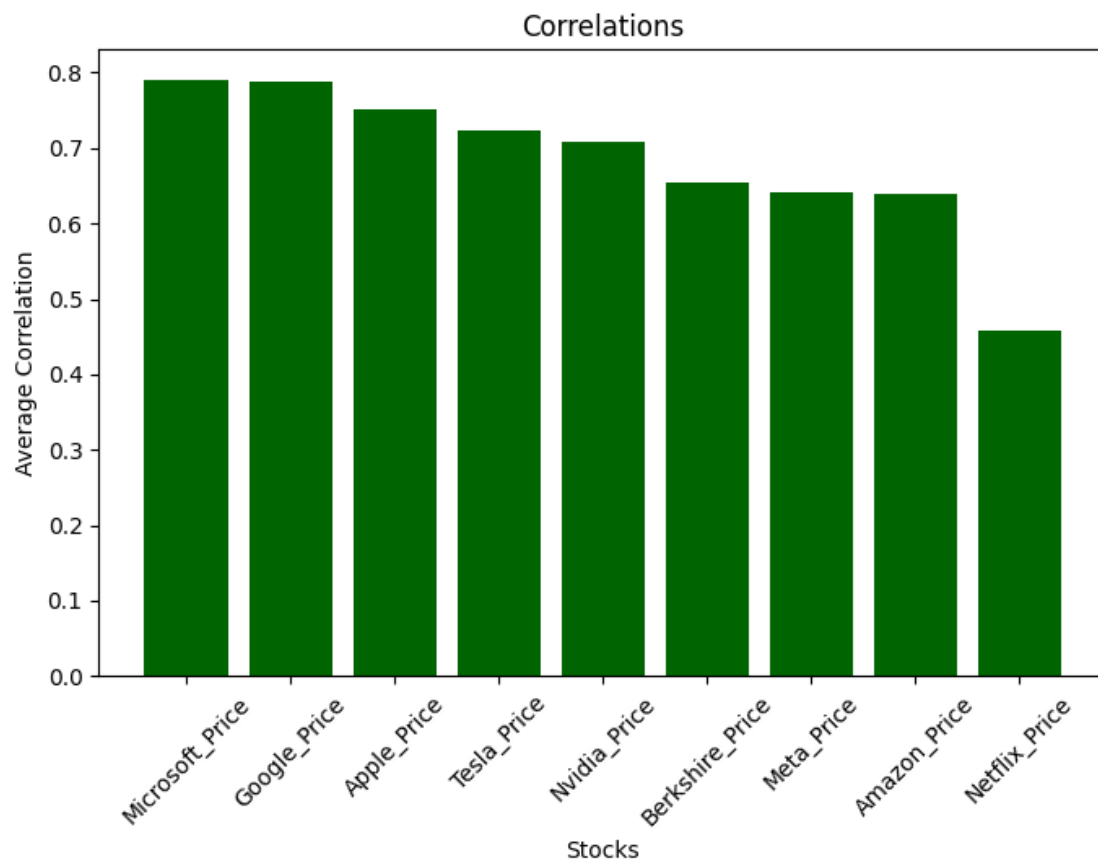
```
[18]:           Average Correlation
Stocks
Microsoft_Price      0.790848
Google_Price         0.787895
```


Apple_Price	0.751222
Tesla_Price	0.723545
Nvidia_Price	0.709145
Berkshire_Price	0.654055
Meta_Price	0.640594
Amazon_Price	0.639962
Netflix_Price	0.457517

```
[19]: plt.figure(figsize=(8, 5))
plt.bar(ac_frame.index, ac_frame['Average Correlation'], color='darkgreen')

plt.xticks(ac_frame.index, rotation=45)

plt.xlabel("Stocks")
plt.ylabel("Average Correlation")
plt.title("Correlations")
plt.show()
```



1.0.6 6) Market Visualizations

Line Plots:

Analyzing the movements of the markets over time.

Both the S&P500 and Nasdaq prices have grown tremendously in the timeframe. The S&P500 grew by nearly double and the Nasdaq has more than doubled.


```
[20]: colors = ['navy','gray']

fig, axes = plt.subplots(nrows=len(market_prices.columns), ncols=1,
    figsize=(12,10), sharex=True)

for i, (market,color) in enumerate(zip(market_prices.columns,colors)):
    axes[i].plot(df['Date'], df[market], label=market, color=color)
    axes[i].set_title(f'{market} Over Time')
    axes[i].set_ylabel('Price (USD)')
    axes[i].legend()

plt.xticks(rotation=45)
plt.xlabel('Date')

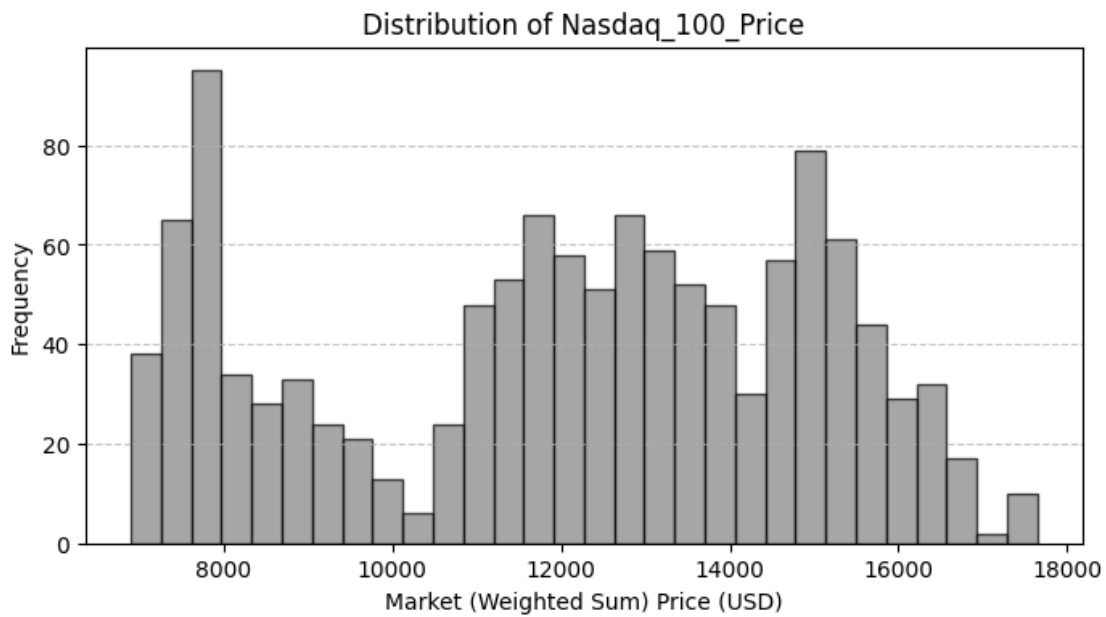

plt.tight_layout()
plt.show()
```



Histogram:

Check distribution of prices over the dataset period.

```
[21]: for market, color in zip(market_prices.columns, colors):
    plt.figure(figsize=(8, 4))
    plt.hist(market_prices[market], bins=30, color=color, alpha=0.7,
             edgecolor='black')
    plt.title(f'Distribution of {market}')
    plt.xlabel('Market (Weighted Sum) Price (USD)')
    plt.ylabel('Frequency')
    plt.grid(axis='y', linestyle='--', alpha=0.7)
    plt.show()
```



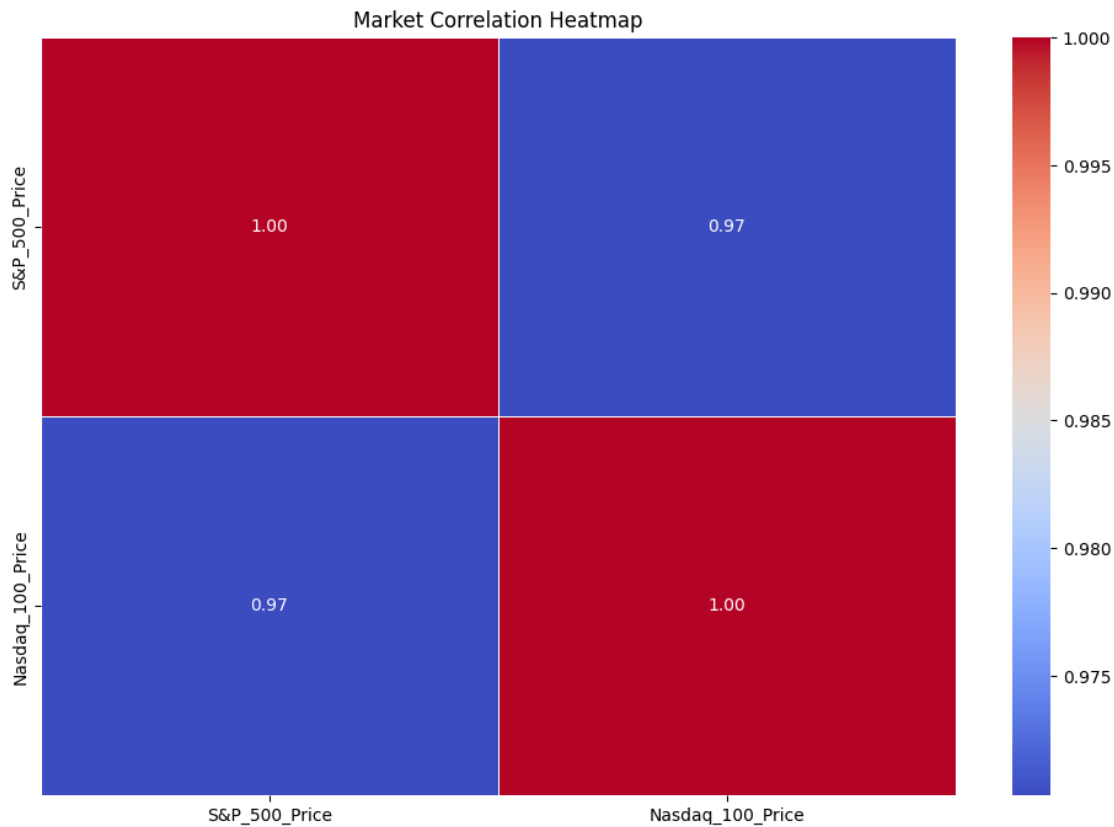
Heat Map:

Check for correlations of market prices between Nasdaq and S&P500.

Strong positive correlation between the them (0.985).

```
[22]: corr_matrix = market_prices.corr()

plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title('Market Correlation Heatmap')
plt.show()
```



```
[23]: average_correlations=[]
for column in corr_matrix.columns:
    average_correlations.append([column,corr_matrix[column].mean()])
ac_frame = pd.DataFrame(average_correlations,columns=['Markets','Average_
Correlation']).sort_values(by='Average Correlation',ascending=False).
set_index('Markets')
ac_frame
```

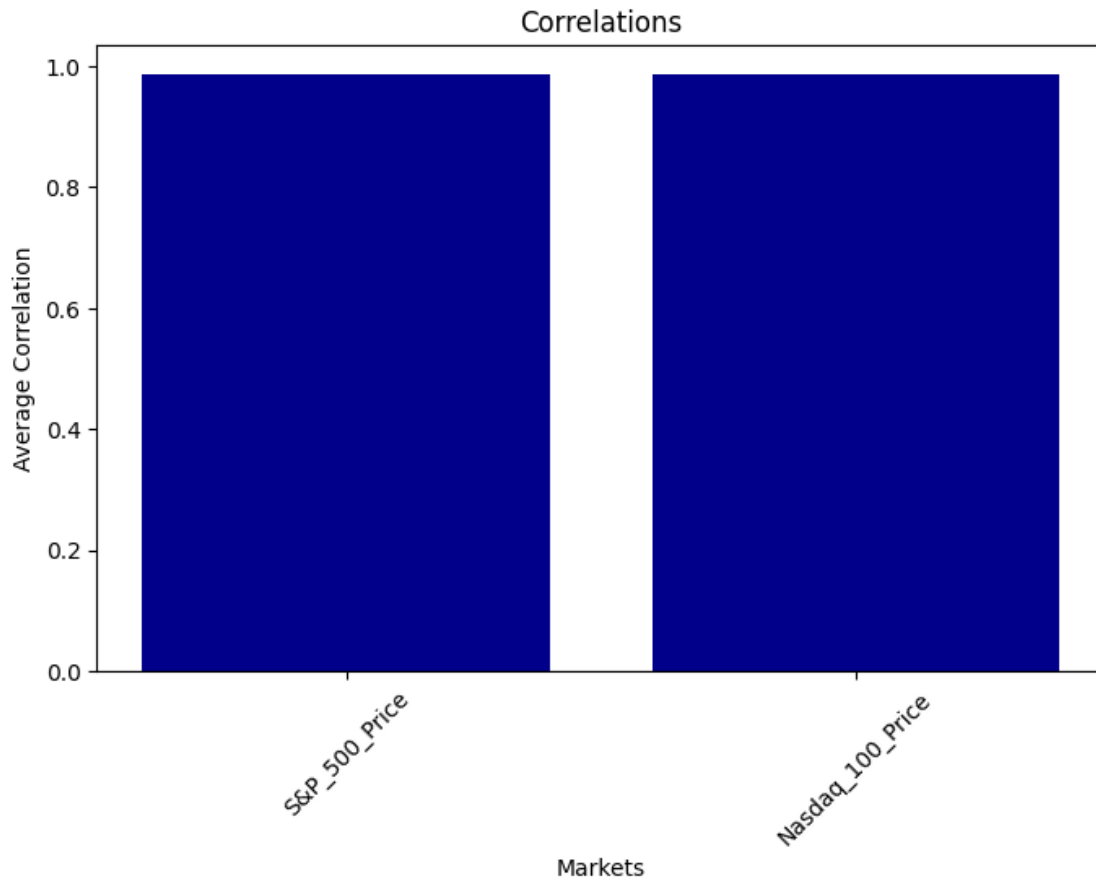
```
[23]:
```

Markets	Average Correlation
S&P_500_Price	0.985157
Nasdaq_100_Price	0.985157

```
[24]: plt.figure(figsize=(8, 5))
plt.bar(ac_frame.index, ac_frame['Average Correlation'], color='darkblue')

plt.xticks(ac_frame.index, rotation=45)

plt.xlabel("Markets")
plt.ylabel("Average Correlation")
plt.title("Correlations")
plt.show()
```



1.0.7 7) Commodity Visualizations

Line Plots:

Analyzing the movements of the commodities over time.

Crude oil price went below zero for a brief moment between 2020 and 2021.

Most of the commodities have not changed in value over time. The main exception here is gold. Gold has risen by 500 USD between 2019 and 2024 according to the data.


```
[25]: colors = sns.color_palette("tab10", len(commodity_prices.columns))

fig, axes = plt.subplots(nrows=len(commodity_prices.columns), ncols=1,
    figsize=(12,10), sharex=True)

for i, (commodity,color) in enumerate(zip(commodity_prices.columns,colors)):
    axes[i].plot(df['Date'], df[commodity], label=commodity, color=color)
    axes[i].set_title(f'{commodity} Over Time')
    axes[i].set_ylabel('Price (USD)')
    axes[i].legend()

plt.xticks(rotation=45)
plt.xlabel('Date')


plt.tight_layout()
plt.show()
```





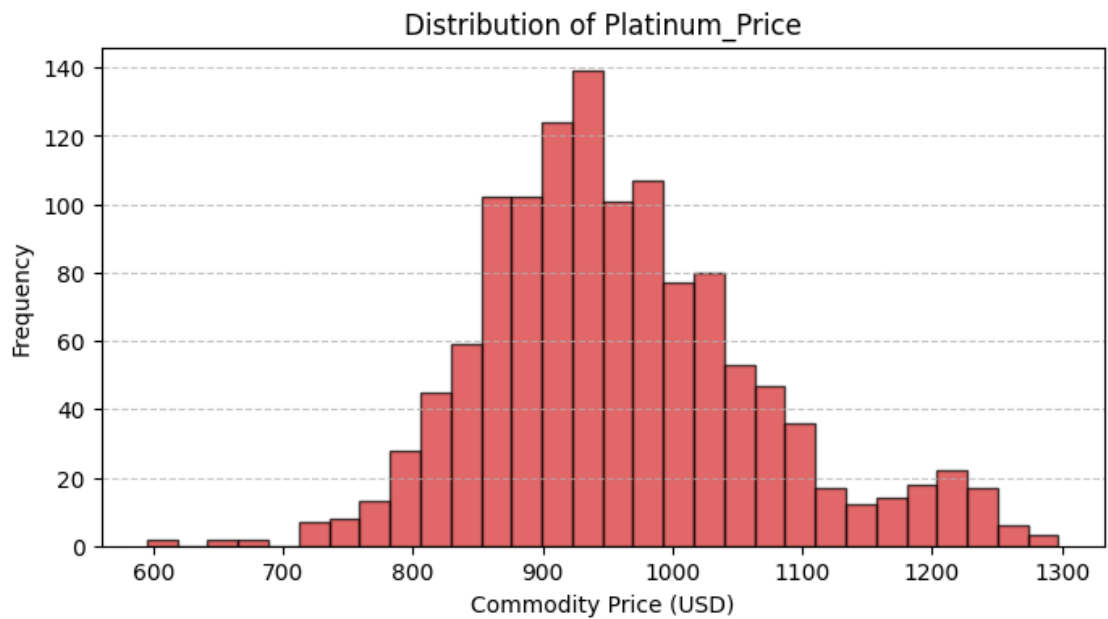
Histogram:

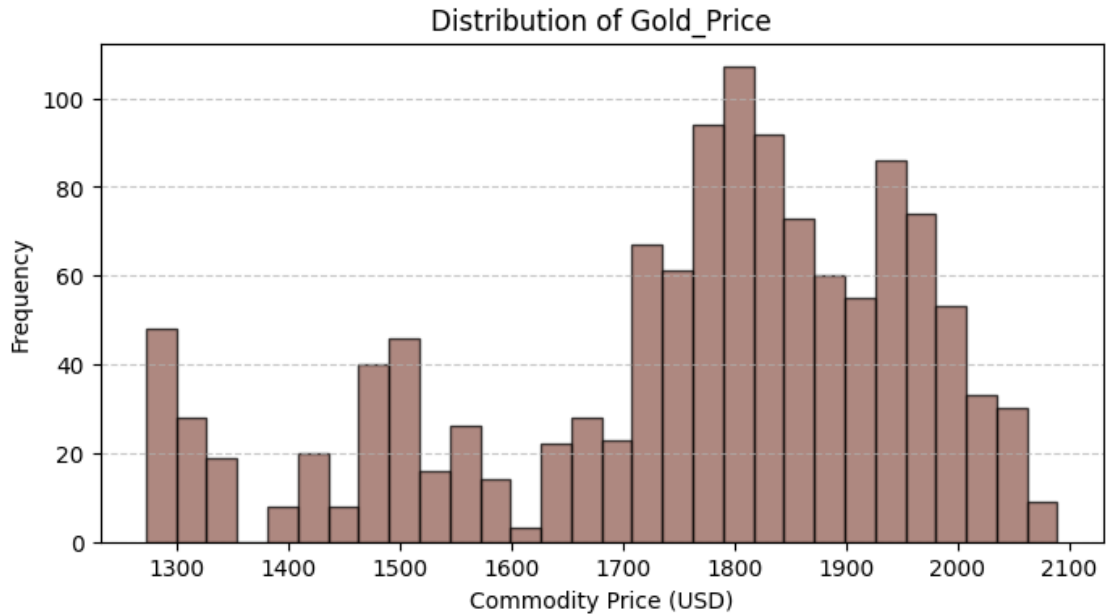
Check distribution of prices over the dataset period.

```
[26]: for commodity, color in zip(commodity_prices.columns, colors):  
    plt.figure(figsize=(8, 4))  
    plt.hist(commodity_prices[commodity], bins=30, color=color, alpha=0.7,  
             edgecolor='black')  
    plt.title(f'Distribution of {commodity}')  
    plt.xlabel('Commodity Price (USD)')  
    plt.ylabel('Frequency')  
    plt.grid(axis='y', linestyle='--', alpha=0.7)  
    plt.show()
```









Heat Map:


Check for correlations of prices between commodities.

Most correlated to other commodities on average is copper (0.70). Its strongest relationships are with crude oil and silver.

Natural gas was the least correlated to the others on average (0.42).

```
[27]: corr_matrix = commodity_prices.corr()

plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title('Commodity Correlation Heatmap')
plt.show()
```



```
[28]: average_correlations=[]
for column in corr_matrix.columns:
    average_correlations.append([column,corr_matrix[column].mean()])
ac_frame = pd.DataFrame(average_correlations,columns=['Commodity','Average_
↪Correlation']).sort_values(by='Average Correlation',ascending=False).
↪set_index('Commodity')
ac_frame
```

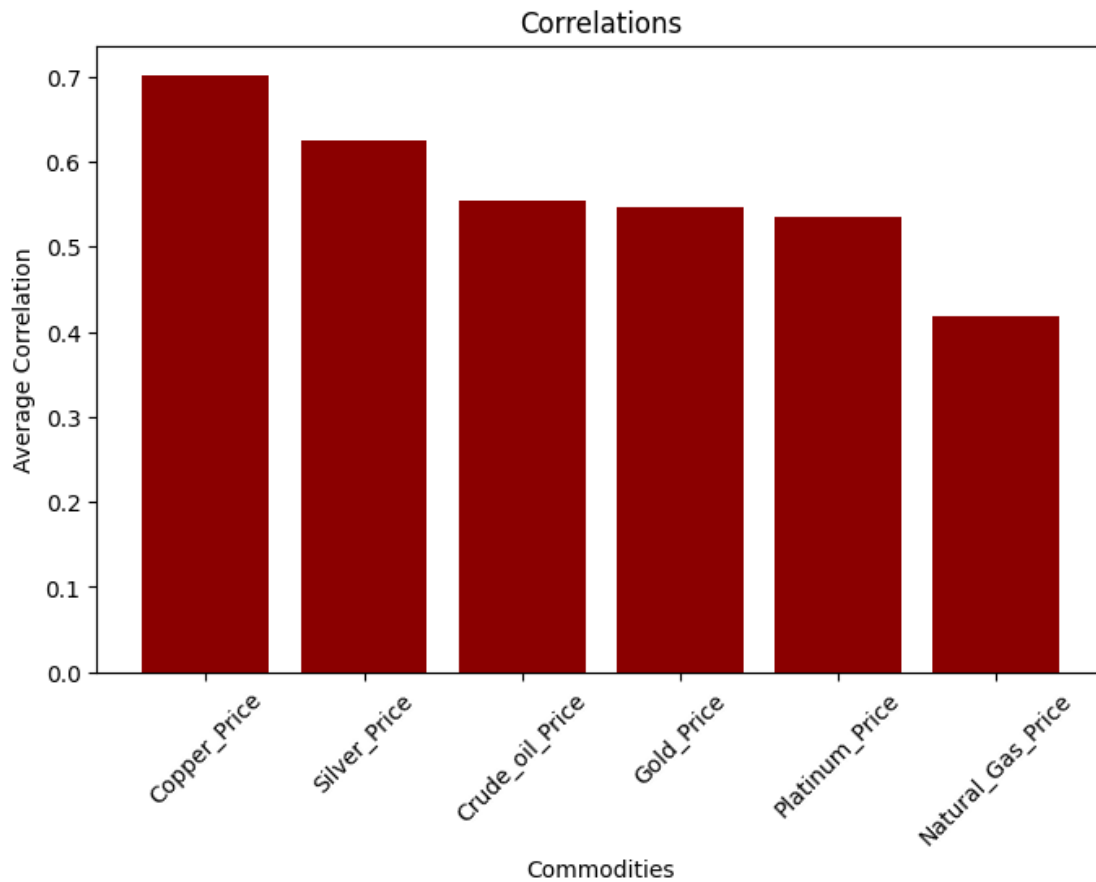
```
[28]:                Average Correlation
Commodity
Copper_Price          0.701322
Silver_Price          0.625721
Crude_oil_Price       0.553617
Gold_Price            0.546527
Platinum_Price        0.534400
Natural_Gas_Price     0.417986
```

```
[29]: plt.figure(figsize=(8, 5))
plt.bar(ac_frame.index, ac_frame['Average Correlation'], color='darkred')

plt.xticks(ac_frame.index, rotation=45)

plt.xlabel("Commodities")
```

```
plt.ylabel("Average Correlation")
plt.title("Correlations")
plt.show()
```



1.0.8 8) Combined Analysis

```
[30]: merged = df[['Apple_Price', 'Tesla_Price', 'Microsoft_Price', 'Google_Price',
                  'Nvidia_Price', 'Berkshire_Price', 'Netflix_Price', 'Amazon_Price',
                  'Meta_Price', 'S&P_500_Price',
                  ↪ 'Nasdaq_100_Price', 'Natural_Gas_Price',
                  'Crude_oil_Price', 'Copper_Price', 'Platinum_Price',
                  ↪ 'Silver_Price',
                  'Gold_Price']]
```

```
[31]: merged
```

```
[31]:
```

	Apple_Price	Tesla_Price	Microsoft_Price	Google_Price	Nvidia_Price	\
0	185.85	187.91	411.22	142.38	661.60	

1	186.86	188.86	403.78	141.16	630.27
2	184.40	187.29	397.58	140.10	615.27
3	188.04	191.59	408.59	151.46	627.74
4	191.73	190.93	409.72	153.51	624.65
...
1238	42.60	20.39	105.67	55.12	37.04
1239	42.73	20.50	105.27	55.30	36.85
1240	43.56	21.15	106.03	56.14	38.25
1241	43.55	21.42	107.22	57.59	37.49
1242	42.81	20.86	105.74	57.07	37.30

	Berkshire_Price	Netflix_Price	Amazon_Price	Meta_Price	S&P_500_Price	\
0	589498.0	564.64	171.81	474.99	4958.61	
1	581600.0	567.51	159.28	394.78	4906.19	
2	578020.0	564.11	155.20	390.14	4848.87	
3	584680.0	562.85	159.00	400.06	4924.97	
4	578800.0	575.79	161.26	401.02	4927.93	
...	
1238	300771.0	347.57	79.41	167.33	2707.88	
1239	302813.0	344.71	80.72	166.38	2706.05	
1240	308810.0	352.19	82.01	170.49	2731.61	
1241	310700.0	355.81	82.94	171.16	2737.70	
1242	312000.0	351.34	81.67	169.25	2724.87	

	Nasdaq_100_Price	Natural_Gas_Price	Crude_oil_Price	Copper_Price	\
0	17642.73	2.079	72.28	3.8215	
1	17344.71	2.050	73.82	3.8535	
2	17137.24	2.100	75.85	3.9060	
3	17476.71	2.077	77.82	3.9110	
4	17596.27	2.490	76.78	3.8790	
...	
1238	6913.13	2.583	52.72	2.8140	
1239	6904.98	2.551	52.64	2.8320	
1240	6997.62	2.662	54.01	2.8400	
1241	7023.52	2.662	53.66	2.8205	
1242	6959.96	2.660	54.56	2.7975	

	Platinum_Price	Silver_Price	Gold_Price
0	901.60	22.796	2053.7
1	922.30	23.236	2071.1
2	932.60	23.169	2067.4
3	931.70	23.225	2050.9
4	938.30	23.134	2034.9
...
1238	802.20	15.809	1318.5
1239	800.80	15.713	1314.2
1240	807.10	15.701	1314.4

1241	821.35	15.836	1319.2
1242	822.50	15.886	1319.3

[1243 rows x 17 columns]

Heat Map:


Check for correlations of prices between commodities, stocks and markets.

It is logical that the markets have the highest correlations (0.76) as they represent the combination of all the prices.

The least correlated are natural gas prices (0.29). This means it is not linearly dependent on the other variables here.

```
[32]: corr_matrix = merged.corr()

plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title('Price Correlation Heatmap')
plt.show()
```



```
[33]: average_correlations=[]
      for column in corr_matrix.columns:
          average_correlations.append([column,corr_matrix[column].mean()])
ac_frame = pd.DataFrame(average_correlations,columns=['Price','Average_
↳Correlation']).sort_values(by='Average Correlation',ascending=False).
↳set_index('Price')
ac_frame
```


```
[33]:
```

	Average Correlation
Price	
Nasdaq_100_Price	0.769442
S&P_500_Price	0.764756
Google_Price	0.751827
Microsoft_Price	0.738718
Apple_Price	0.730277
Tesla_Price	0.718964
Copper_Price	0.696349
Berkshire_Price	0.651488
Silver_Price	0.642931
Gold_Price	0.628075
Nvidia_Price	0.622571
Amazon_Price	0.572071
Meta_Price	0.518160
Crude_oil_Price	0.471140
Platinum_Price	0.450478
Netflix_Price	0.340260
Natural_Gas_Price	0.289678

```
[34]: plt.figure(figsize=(10, 8))
      plt.bar(ac_frame.index, ac_frame['Average Correlation'], color='gold')

      plt.xticks(ac_frame.index, rotation=90)

      plt.xlabel("Everything")
      plt.ylabel("Average Correlation")
      plt.title("Correlations")
      plt.show()
```

- **Nasdaq 100 & S&P 500 (0.77 & 0.76 correlation):**

These indices have a strong correlation with commodities and major tech stocks because they represent the broader market. Their movement often reflects macroeconomic conditions, monetary policy, and investor sentiment, which also drive commodity prices and large-cap stocks.

- **Google & Microsoft (0.75 & 0.74 correlation):**

These two tech giants show strong correlation with commodities and the market, likely due to their central role in the economy. Their revenues are tied to advertising, cloud computing, and enterprise software, all of which depend on economic growth and investment cycles.

- **Copper Price (0.70 correlation):**

Copper is an industrial metal widely used in construction and electronics. Its strong correlation with stock prices suggests that investors see copper demand as an indicator of economic expansion, which also benefits companies in tech and manufacturing.

- **Gold Price (0.63 correlation):**

While gold is traditionally a safe-haven asset, its correlation with equities suggests it may have been moving in response to inflation expectations rather than acting as a hedge against market declines.